

# MCDA5580 Assignment 5

## Team Member

- Hemalatha Srinivasan A00452621
- Ajay Jain A00455849
- Kin Wa, Chan A00467755

# Table of Content

Executive Summary.....	3
Introduction .....	4
Objectives .....	4
Data Analysis.....	5
Data Extraction.....	5
Data cleaning .....	7
Design/ Methodology/ Approach .....	11
Overview .....	11
Models Used in Analysis .....	11
Training/ testing data sets .....	11
Sampling Method.....	11
Mean Absolute Percentage Error (MAPE) .....	11
Linear Regression Analysis .....	12
Neural Network Analysis.....	15
Support Vector Machine (SVM) Analysis .....	24
Holt Winters Analysis.....	35
ARIMA Analysis .....	36
Random Forest Analysis.....	38
Conclusion.....	41
Comparison of all models used.....	41
Recommendation and next step.....	43
Appendix .....	44
SQL Script for Data Extraction.....	44
R Script for Data Analysis .....	44
Linear Regression Analysis .....	44
Neural Network Analysis.....	52

SVM Analysis .....	66
HoltWinters, ARIMA and Random Forest Analysis .....	78
Reference/ Citation.....	86

## **Executive Summary**

This report examines recent product sales data obtained from a retail store to identify potential strategies for improving sales performance. The study focuses on analyzing the sales quantity of the top consumed products and predict the next day sales with previous 7 days through using time series prediction methods. It is implemented by extracting daily sales from database and transposed to dataset with the each record contains the sales of 8 continuous days.

The main objective of this study was to find future trends from past data, forecasting future values, identifying patterns and trends, optimizing product stocking, and improving decision-making. The use of these predictive models can help businesses and organizations with inventory planning for the future and make informed decisions to optimize their operations and maximize profits.

The tested methods in this study include Linear Regression Analysis, Neural Network Analysis, Support Vector Machine (SVM) Analysis, HoltWinters Analysis, Autoregressive Integrated Moving Average (ARIMA) Analysis, and Random Forest Analysis. The best model was found to be the averaged neural network and ARIMA models with an error rate of ~30%.

The report recommends using the averaged neural network and ARIMA models for future prediction of sales quantity and maintaining inventory levels at the store appropriately. The report also recommends keeping track of any deviation in the model prediction vs actuals to adjust appropriately and update the model using the latest data.

# **Introduction**

The purpose of this data analysis report is to examine recent product sales data obtained from a retail store and identify potential strategies for improving sales performance. Specifically, this study will focus on analyzing the sales quantity of the top consumed products. To accomplish this, we will employ time series prediction methods in forecasting, using the quantity sold in the past seven days to predict the next day's quantity. Our analysis will utilize statistical, neural network, and support vector regression techniques, and we will compare the performance of these models.

# **Objectives**

The main objective of time series prediction is to find the future trends from the past data. The goal is to find a model which can accurately predict future value. The prediction will help in satisfying below purpose:

1. Forecasting future values: This is the major goal in using Time series predictions. Knowing the future results can increase business efficiency in decision making.
2. Identifying patterns and trends: We can find the patterns and trends using the past value which can be used for predicting the future trends and patterns.
3. Optimizing product stocking: This can help in predicting the best performing product. Using the obtained results, we can predict inventory levels for various products.
4. Improving decision-making: This will help in allocating the budget for the products involved in the analysis.

This can help businesses and organizations with inventory planning for the future and make informed decisions to optimize their operations and maximize their profits.

# Data Analysis

## Data Extraction

The sales data with the following fields between 2/1/2015 and 14/9/2015 are gathered for analysis (total 1328 records)

Field Name	Data Type	Description
date	Date	Sales Date
item_sk	String	Item No
quantity	Integer	Total no of sales the date

The quantities are aggregated by product and top 6 product are ranked out:

item_sk	total_quantity
11740941	90239
11740923	61482
11741274	28261
11636550	23839
11629829	22338
11743201	20716

Top 1<sup>st</sup>, 3<sup>rd</sup>, 5<sup>th</sup> (11740941, 11741274, 11629829) are selected for study.

For each item, the records are transposed for further analysis. Example:

Original record

date	item_sk	quantity
2/1/2015	11740941	430
3/1/2015	11740941	433
4/1/2015	11740941	305
5/1/2015	11740941	469
6/1/2015	11740941	533
7/1/2015	11740941	273
8/1/2015	11740941	279
9/1/2015	11740941	292
10/1/2015	11740941	399
11/1/2015	11740941	362

12/1/2015	11740941	315
13/1/2015	11740941	579
14/1/2015	11740941	364
15/1/2015	11740941	388
16/1/2015	11740941	383
17/1/2015	11740941	393

Transposed record

V1	V2	V3	V4	V5	V6	V7	V8
430	433	305	469	533	273	279	292
399	362	315	579	364	388	383	393

The study is going to use the quantities in previous 7 days (V1 – V7) to predict the quantity in the next day (V8).

Total 31 sets of transposed records are obtained for each product.

## Data cleaning

Box plot analysis is used to figure out the outliers. Extreme outliers are removed:

- For 1<sup>st</sup> top item, removed the item with quantity = 3
- For 3<sup>rd</sup> top item, removed the item with quantity = 1
- For 1<sup>5th</sup> top item, removed the item with quantity = 603

Box plot for 1<sup>st</sup> top item

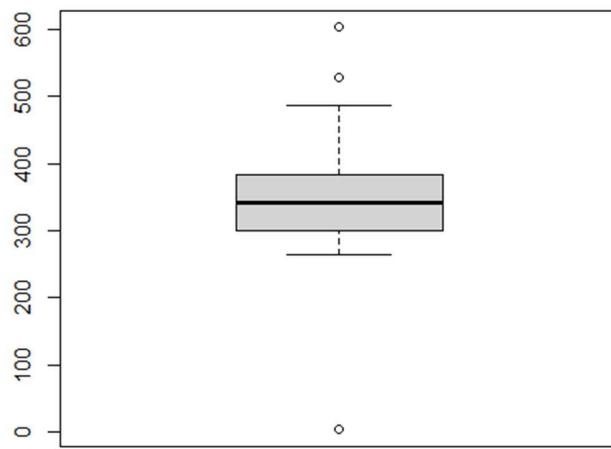


Figure 1

Box plot for 3<sup>rd</sup> top item

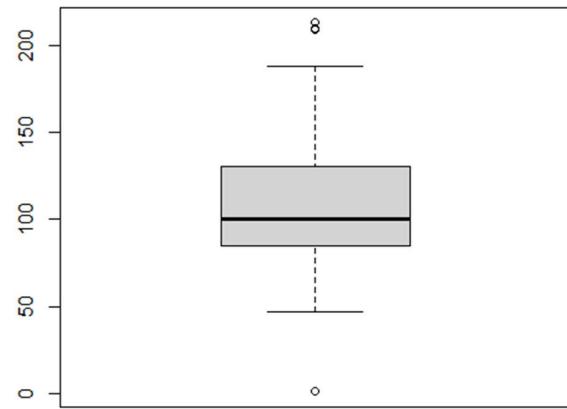


Figure 2

Box plot for 5<sup>th</sup> top item

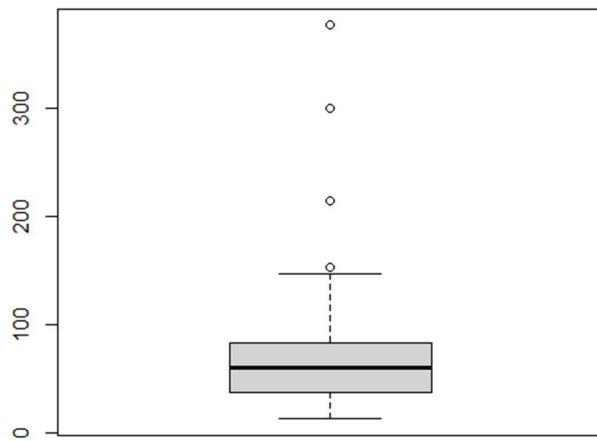


Figure 3

Plot as weekly time series (for three products respectively)

We can observe there are seasonal patterns, however the overall trends are not obvious.

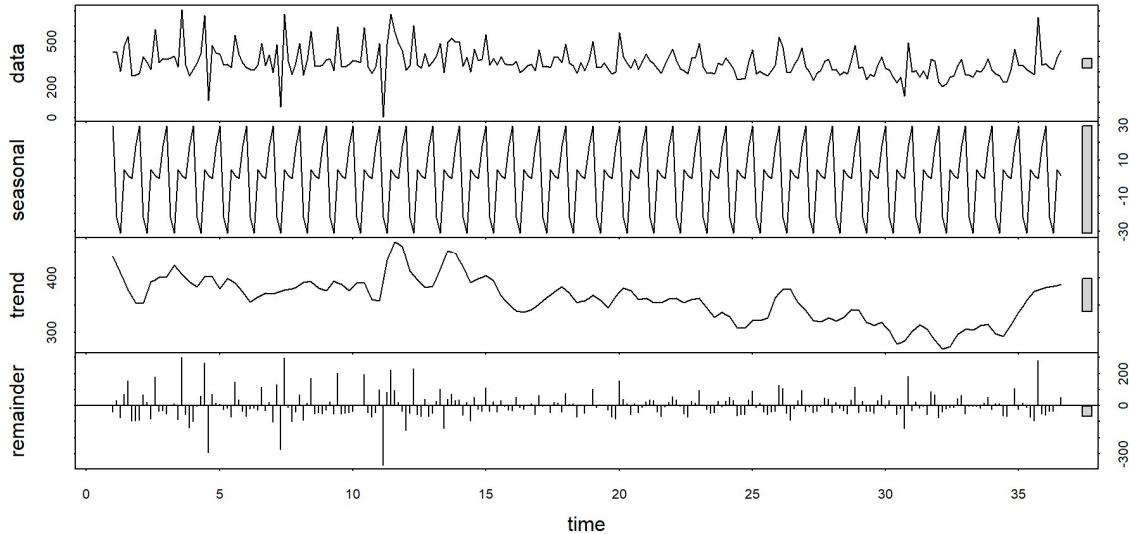


Figure 4

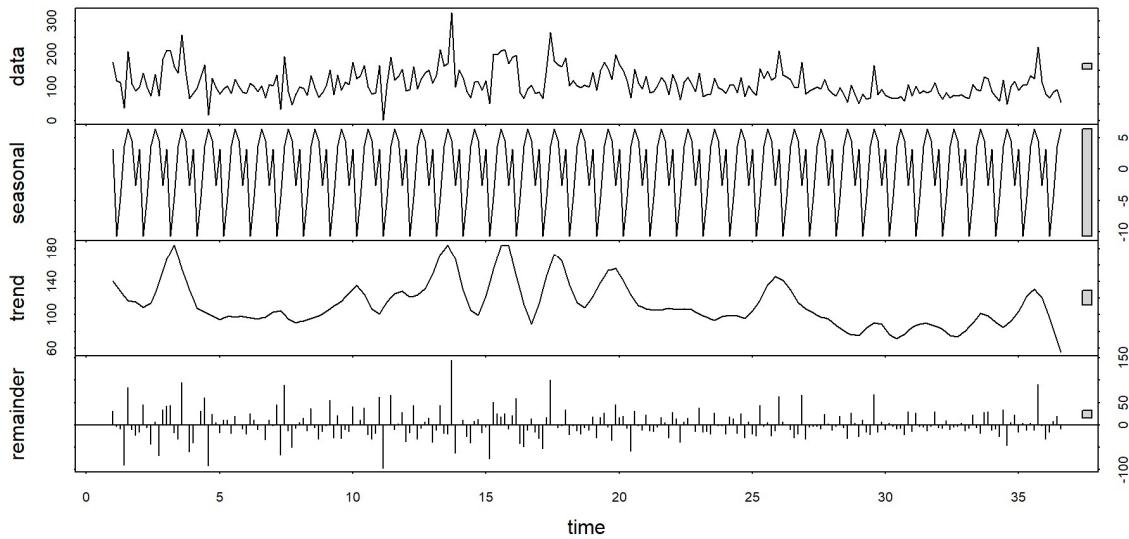


Figure 5

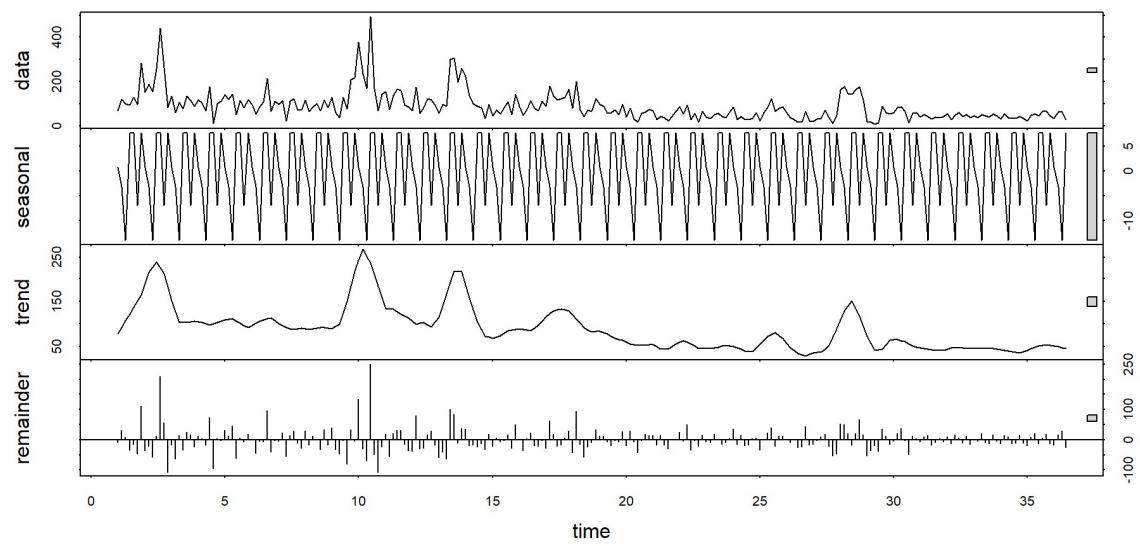


Figure 6

# **Design/ Methodology/ Approach**

## **Overview**

### **Models Used in Analysis**

The following methods are tested in this study:

- Linear Regression Analysis
- Neural Network Analysis
- Support Vector Machine (SVM) Analysis
- HoltWinters Analysis
- Autoregressive Integrated Moving Average (ARIMA) Analysis
- Random Forest Analysis

Each model has its strengths and weaknesses and speciality on specific problems. Through the experiments with these models, we could better understand the nature of the data. In general speaking, Linear regression is simple and interpretable, neural networks and SVMs can handle complex relationships between variables, Holt-Winters and ARIMA are effective for time-series data, and Random Forest can handle both categorical and continuous data.

Caret Package (short for Classification And REgression Training) in R is used to implement these models. The package contains tools for data splitting, pre-processing, feature selection, model tuning using resampling, and variable importance estimation, all helps to streamline the process of creating predictive models.

### **Training/ testing data sets**

To evaluate the performance of models, the studies split the data into training and testing datasets in an 80%/20% ratio.

### **Sampling Method**

Repeated cross-validation (repeatedCV) is used to avoid overfitting and improve the generalizability of the model. Cross-validation is a technique used to evaluate the performance of a model by splitting the data into training and validation sets. Repeated cross-validation takes this a step further by repeating the process multiple times with different random splits of the data.

### **Mean Absolute Percentage Error (MAPE)**

MAPE is used to evaluate the performance of predictions. It is calculated as the average absolute percent error for each time-period minus actual values divided by actual values. The lower value of MAPE means the model has a better accuracy in prediction.

## Linear Regression Analysis

Linear regression is a statistical method that aims to discover linear relationship between predictors and dependent variable, it is used to make predictions with continuous outcomes. It is easy to implement and interpret. However, it is not capable of dealing with nonlinear relationship between variables.

Result summary

1<sup>st</sup> Item

```
> summary(glmFitTime)

Call:
NULL

Deviance Residuals:
    Min      1Q  Median      3Q      Max
-164.620 -42.292 -0.005   45.095  131.736

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 351.667    15.027  23.402 1.8e-15 ***
V1          3.385     19.099   0.177   0.861    
V2          3.888     17.015   0.228   0.822    
V3          14.645    16.626   0.881   0.389    
V4          -5.533    17.299  -0.320   0.753    
V5          -10.380   17.526  -0.592   0.561    
V6          10.828    17.289   0.626   0.539    
V7          27.416    19.023   1.441   0.166    
---
Signif. codes:  0  ***  0.001  **  0.01  *  0.05  .  0.1  '  1

(Dispersion parameter for gaussian family taken to be 6097.067)

Null deviance: 160218  on 26  degrees of freedom
Residual deviance: 115844  on 19  degrees of freedom
AIC: 320.46
```

3<sup>rd</sup> Top item

```

> summary(glmFitTime)

Call:
NULL

Deviance Residuals:
    Min      1Q  Median      3Q     Max 
-49.443 -19.106 -9.702  15.997  75.577 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 109.704    6.642   16.518 9.99e-13 ***
V1           7.858    7.633   1.029   0.316    
V2           5.752    8.669   0.663   0.515    
V3          -4.781    8.060   -0.593   0.560    
V4          -14.374   10.381  -1.385   0.182    
V5          -4.475    7.778   -0.575   0.572    
V6           7.554    9.430   0.801   0.433    
V7          26.064    9.741   2.676   0.015 *  
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for gaussian family taken to be 1191.014)

Null deviance: 38446 on 26 degrees of freedom
Residual deviance: 22629 on 19 degrees of freedom
AIC: 276.36

Number of Fisher Scoring iterations: 2

5th Top item

```

```

> summary(glmFitTime)

Call:
NULL

Deviance Residuals:
    Min      1Q  Median      3Q     Max 
-58.661 -22.065 -13.948   3.842 189.033 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 78.074     11.435   6.828 1.62e-06 ***
V1          -1.602     21.331  -0.075   0.941    
V2          26.200     20.896   1.254   0.225    
V3         -20.331     22.031  -0.923   0.368    
V4          -7.051     29.779  -0.237   0.815    
V5          -3.294     27.560  -0.120   0.906    
V6          24.181     19.625   1.232   0.233    
V7          12.770     22.087   0.578   0.570    
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for gaussian family taken to be 3530.286)

Null deviance: 101130  on 26  degrees of freedom
Residual deviance: 67075  on 19  degrees of freedom
AIC: 305.7

Number of Fisher Scoring iterations: 2

```

We can observe the P value of the predictors are not statistically significant ( $>0.05$ ) except the V7 predictor in 3<sup>rd</sup> top item. It shows linear regression is not an appropriate model for the data set.

#### MAPE Summary

MAPE trends to be high especially for the training set. It may be another evidence that shows linear regression is not a suitable model.

Item	Training	Testing
1st Top Item	16.91%	23.78%
3rd Top Item	23.43%	28.53%
5th Top Item	44.80%	172.90%
Average	28.38%	75.07%

## **Neural Network Analysis**

Neural networks are a type of machine learning model that is inspired by the structure and function of the human brain. They consist of interconnected nodes or neurons that work together to learn patterns in the input data and make predictions or decisions based on that learning.

The advantages include being able to model nonlinear relationships, high accuracy, and high fault tolerance to noisy data. However, it had the drawback of black box nature (relationship learned is difficult to explain), high computational power requirement, and require large amount of data for training.

There are two main parameters control capacity of Neural networks model:

1. Number of nodes: Increasing the number of nodes in a single layer can improve the model's ability to capture complex relationships between the input and output.
2. Number of layers: Increasing the number of layers not only increases the capacity of the model, but also enables it to capture nonlinear relationship.

Increasing these parameters could improve the ability to recognize features but it causes overfitting problem.

Neural networks with different configurations are tested in this analysis:

1. Neural network with single layer, number of nodes = 15
2. Neural network with 3 layers, 5 nodes in each layer
3. Averaged Neural Network – It trains multiple neural network models and then averaging their predictions to produce the final output. The main advantage is to overcome the overfitting problem of normal neural network but resulting more computationally expensive.

## Result summary

Neural network with single layer, number of nodes = 15

1<sup>st</sup> Item

Neural network Plot

It demonstrates the single layer neural network architecture and its parameters

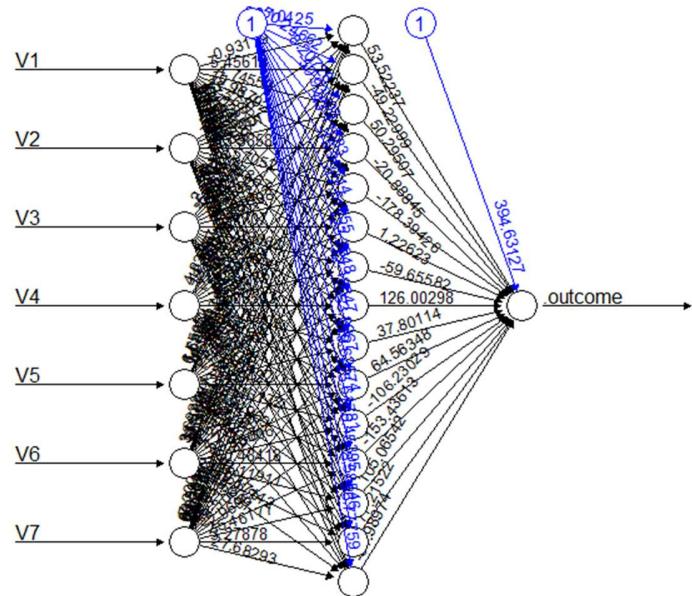


Figure 7

## Model summary

It shows the model configuration used.

```
> summary(nnFitTime)
   Length Class      Mode
call        7  -none-    call
response     27  -none-   numeric
covariate   189  -none-   numeric |
model.list    2  -none-    list
err.fct       1  -none-   function
act.fct       1  -none-   function
linear.output  1  -none-   logical
data         8  data.frame list
exclude       0  -none-   NULL
net.result     1  -none-    list
weights       1  -none-    list
generalized.weights  1  -none-    list
startweights    1  -none-    list
result.matrix  139  -none-   numeric
xNames        7  -none-   character
problemType    1  -none-   character
tuneValue      3  data.frame list
obsLevels      1  -none-   logical
param        ...  -none-    list
```

3<sup>rd</sup> Top item

### Neural network Plot

It demonstrates the single layer neural network architecture and its parameters.

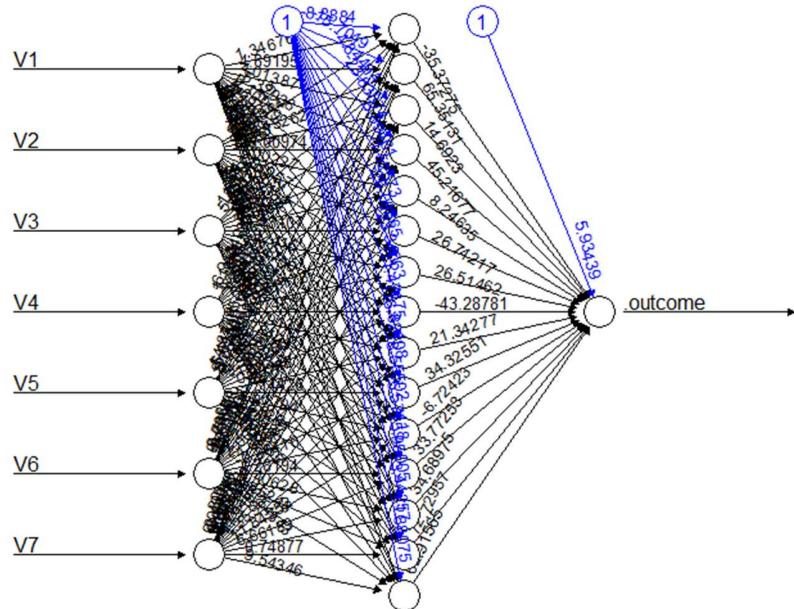


Figure 8

### Model summary

It shows the model configuration used.

```
> summary(nnFitTime)
   Length Class      Mode
call          8  -none-    call
response      27  -none-   numeric
covariate     189  -none-   numeric
model.list    2   -none-   list
err.fct       1   -none-   function
act.fct       1   -none-   function
linear.output 1   -none-   logical
data          8   data.frame list
exclude        0   -none-   NULL
net.result     1   -none-   list
weights        1   -none-   list
generalized.weights 1   -none-   list
startweights   1   -none-   list
result.matrix 139  -none-   numeric
xNames         7   -none-   character
problemType    1   -none-   character
tuneValue      3   data.frame list
obsLevels      1   -none-   logical
param          4   -none-   list
```

5<sup>th</sup> Top item

### Neural network Plot

It demonstrates the single layer neural network architecture and its parameters

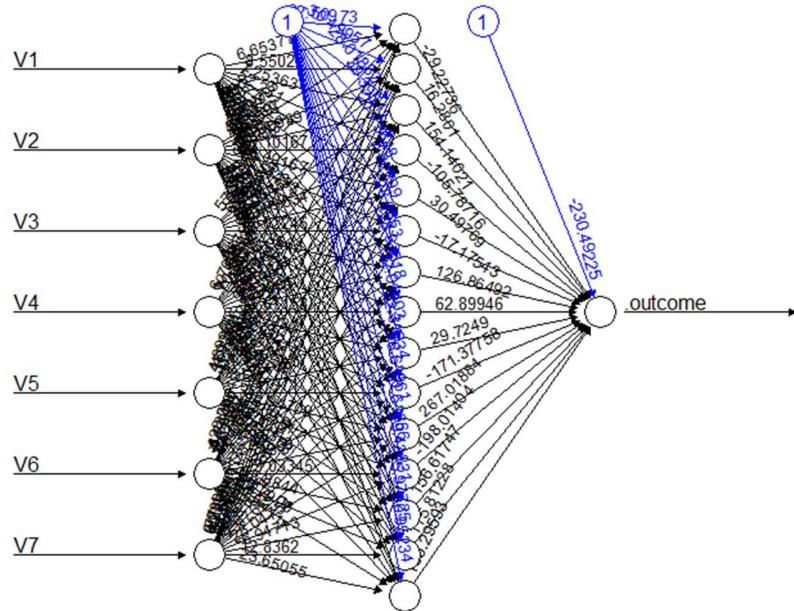


Figure 9

### Model summary

It shows the model configuration used.

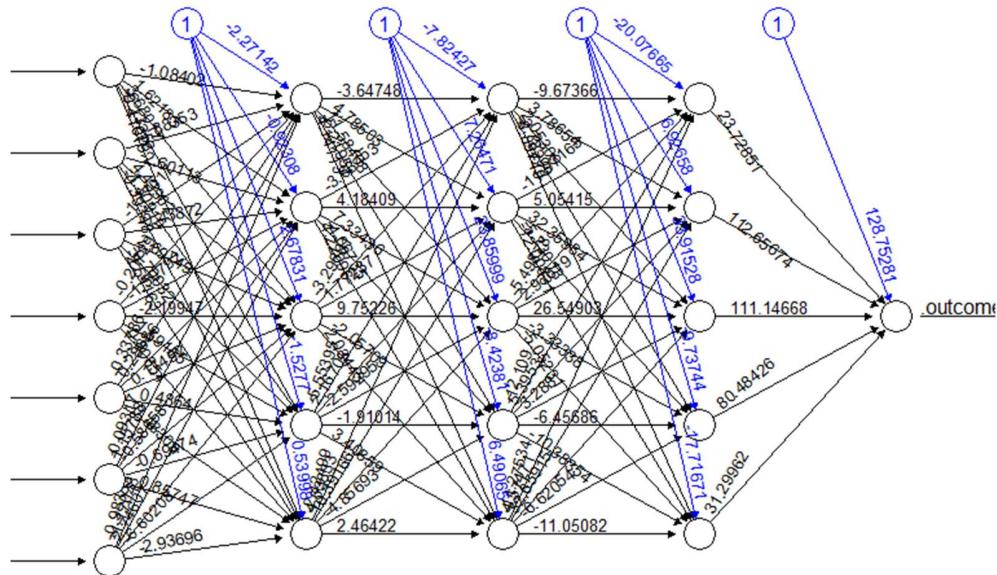
```
> summary(nnFitTime)
      Length Class      Mode
call        7 -none-    call
response    27 -none- numeric
covariate   189 -none- numeric
model.list  2 -none-  list
err.fct     1 -none- function
act.fct     1 -none- function
linear.output 1 -none- logical
data        8 data.frame list
exclude     0 -none-  NULL
net.result  1 -none-  list
weights     1 -none-  list
generalized.weights 1 -none-  list
startweights 1 -none-  list
result.matrix 139 -none- numeric
xNames      7 -none- character
problemType 1 -none- character
tuneValue    3 data.frame list
obsLevels   1 -none- logical
param       3 -none-  list
```

Neural network with 3 layers, 5 nodes in each layer

1<sup>st</sup> Item

Neural network Plot

It demonstrates the 3 layers neural network architecture and their parameters.



Error: 86154.333339 Steps: 366

Figure 10

Model summary

It shows the model configuration used.

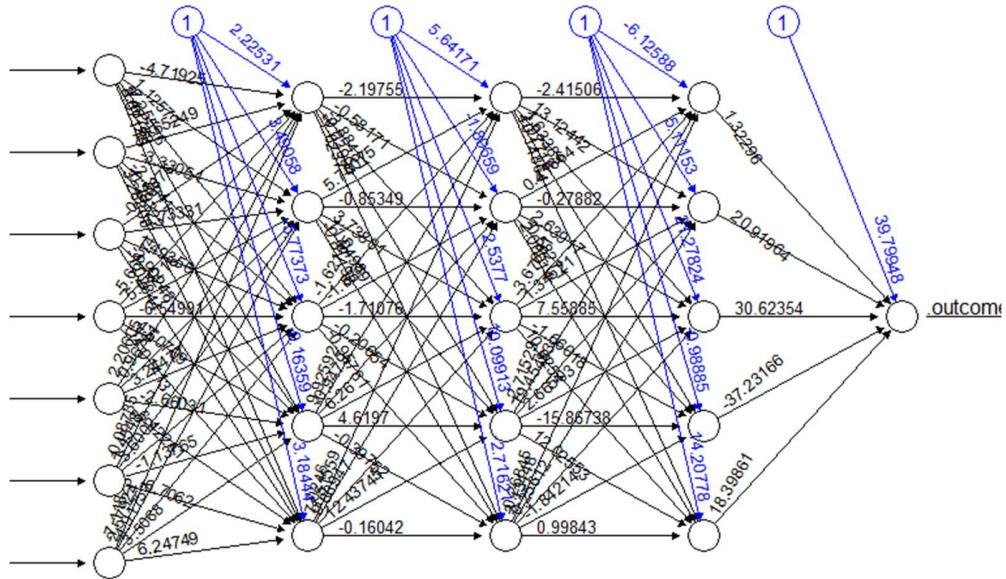
```
> summary(nnFitTime)
   .....
```

	Length	Class	Mode
call	7	-none-	call
response	27	-none-	numeric
covariate	189	-none-	numeric
model.list	2	-none-	list
err.fct	1	-none-	function
act.fct	1	-none-	function
linear.output	1	-none-	logical
data	8	data.frame	list
exclude	0	-none-	NULL
net.result	1	-none-	list
weights	1	-none-	list
generalized.weights	1	-none-	list
startweights	1	-none-	list
result.matrix	109	-none-	numeric
xNames	7	-none-	character
problemType	1	-none-	character
tuneValue	3	data.frame	list
obsLevels	1	-none-	logical
param	3	-none-	list

3<sup>rd</sup> Top item

### Neural network Plot

It demonstrates the 3 layers neural network architecture and their parameters.



Error: 19511.591587 Steps: 473

Figure 11

### Model summary

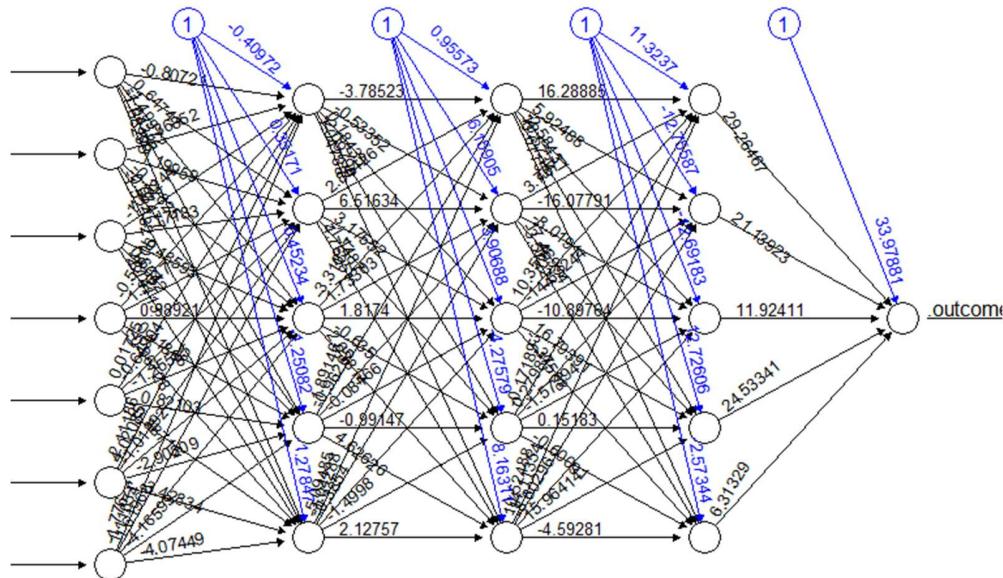
It shows the model configuration used.

```
> summary(nnFitTime)
   Length Class      Mode
call       7 -none-   call
response   27 -none- numeric
covariate 189 -none- numeric
model.list  2 -none-  list
err.fct    1 -none- function
act.fct    1 -none- function
linear.output 1 -none- logical
data       8 data.frame list
exclude    0 -none-  NULL
net.result 1 -none-  list
weights    1 -none-  list
generalized.weights 1 -none-  list
startweights 1 -none-  list
result.matrix 109 -none- numeric
xNames     7 -none- character
problemType 1 -none- character
tuneValue   3 data.frame list
obsLevels   1 -none- logical
param      3 -none-  list
```

5<sup>th</sup> Top item

### Neural network Plot

It demonstrates the 3 layers neural network architecture and their parameters.



Error: 94267.334411 Steps: 51

Figure 12

### Model summary

```
> summary(nnFitTime)
```

	Length	Class	Mode
call	7	-none-	call
response	27	-none-	numeric
covariate	189	-none-	numeric
model.list	2	-none-	list
err.fct	1	-none-	function
act.fct	1	-none-	function
linear.output	1	-none-	logical
data	8	data.frame	list
exclude	0	-none-	NULL
net.result	1	-none-	list
weights	1	-none-	list
generalized.weights	1	-none-	list
startweights	1	-none-	list
result.matrix	109	-none-	numeric
xNames	7	-none-	character
problemType	1	-none-	character
tuneValue	3	data.frame	list
obsLevels	1	-none-	logical
param	3	-none-	list

## Averaged Neural Network

1<sup>st</sup> Item

Model summary

It shows the model configuration used.

```
> summary(nnFitTime)
      Length Class    Mode
model      5   -none-  list
repeats    1   -none-  numeric
bag        1   -none-  logical
seeds      5   -none-  numeric
names      7   -none-  character
terms      3   terms   call
coefnames  7   -none-  character
xlevels    0   -none-  list
xNames     7   -none-  character
problemType 1   -none-  character
tuneValue   3   data.frame list
obsLevels  1   -none-  logical
param      4   -none-  list
...       ...  ...    ...
```

3<sup>rd</sup> Top item

Model summary

It shows the model configuration used.

```
> summary(nnFitTime)
      Length Class    Mode
model      5   -none-  list
repeats    1   -none-  numeric
bag        1   -none-  logical
seeds      5   -none-  numeric
names      7   -none-  character
terms      3   terms   call
coefnames  7   -none-  character
xlevels    0   -none-  list
xNames     7   -none-  character
problemType 1   -none-  character
tuneValue   3   data.frame list
obsLevels  1   -none-  logical
param      4   -none-  list
...       ...  ...    ...
```

5<sup>th</sup> Top item

Model summary

It shows the model configuration used.

```

> summary(nnFitTime)
      Length Class      Mode
model      5    -none-   list
repeats    1    -none-   numeric
bag        1    -none-   logical
seeds      5    -none-   numeric
names      7    -none-   character
terms      3    terms    call
coefnames  7    -none-   character
xlevels    0    -none-   list
xNames     7    -none-   character
problemType 1    -none-   character
tuneValue   3    data.frame list
obsLevels   1    -none-   logical
param      4    -none-   list

```

### MAPE Summary

We can see the 1 layer neural had a very lowest MAPE value in training set, but it is unstable that also shows the highest value among all configurations in its testing set result. The 3 layers neural network and averaged neural network had similar performance between their training and testing sets. However, the averaged neural network shows lower MAPEs compare with those of 3 layers neural network, therefore it is the best performed model.

### Neural network with single layer, number of nodes = 15

Item	Training	Testing
1st Top Item	10.35%	27.76%
3rd Top Item	9.35%	68.94%
5th Top Item	17.18%	97.36%
Average	12.29%	64.69%

### Neural network with 3 layers, 5 nodes in each layer

Item	Training	Testing
1st Top Item	19.91%	16.16%
3rd Top Item	21.31%	58.21%
5th Top Item	93.81%	68.23%
Average	45.01%	47.53%

### Averaged Neural Network

Configuration	Item	Training	Testing
size = 3 (Autotune)	1st Top Item	13.90%	24.74%
size = 1 (Autotune)	3rd Top Item	22.40%	53.19%
size = 1 (Autotune)	5th Top Item	55.48%	25.39%
	Average	30.59%	34.44%

## Support Vector Machine (SVM) Analysis

Support Vector Machine (SVM), a form of supervised learning algorithm, can perform classification or regression tasks. The central concept is to identify a hyperplane that effectively separates the distinct classes present in the training data. SVM shared some similarity with neural network. They could be able to handle high dimensional data and the result models are difficult to explain. However, it requires lesser computation power, lesser sensitive overfitting, more suitable for smaller and less complex data sets.

It allows use of different kernels in modelling. The kernel is a function measuring the similarity of two data points. Changing kernel function allows us to categorize different data patterns.

SVM with different kernels will be tested in the analysis:

1. Radial kernel – It is the most used kernel in SVM which is a flexible method that can capture complex nonlinear relationships between the predictor variables and the response variable.
2. Polynomial kernel – It could also model nonlinear relationships but for less complex data only.
3. Linear kernel – It is suitable when the relationship between the predictor variables and the response variable is linear.
4. Sigmoid kernel - It is like a two-layer perceptron model of the neural network.

## Result Summary

### Radial kernel

#### 1<sup>st</sup> Top item

#### Model summary

It shows the model configuration used and some performance measurements.

#### Support Vector Machines with Radial Basis Function Kernel

27 samples  
7 predictor

Pre-processing: centered (7), scaled (7)

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 24, 25, 24, 25, 24, 24, ...

Resampling results across tuning parameters:

C	RMSE	Rsquared	MAE
0.25	64.08148	0.7975307	53.88688
0.50	62.42463	0.7610670	51.97124
1.00	65.37114	0.7222749	55.19914
2.00	71.44978	0.6975062	60.53483
4.00	78.35225	0.6853026	66.24954
8.00	82.13271	0.6843033	69.74372
16.00	82.05953	0.6852691	69.68440
32.00	82.05953	0.6852691	69.68440
64.00	82.05953	0.6852691	69.68440
128.00	82.05953	0.6852691	69.68440

Tuning parameter 'sigma' was held constant at a value of 0.2024405

RMSE was used to select the optimal model using the smallest value.

The final values used for the model were sigma = 0.2024405 and C = 0.5.

> summary(svmFitTime)

Length Class Mode  
1 ksvm S4

3<sup>rd</sup> Top item

#### Model summary

It shows the model configuration used and some performance measurements.

```
> svmFitTime
Support Vector Machines with Radial Basis Function Kernel

27 samples
 7 predictor

Pre-processing: centered (7), scaled (7)
Resampling: Cross-Validated (10 fold, repeated 5 times)
Summary of sample sizes: 25, 24, 24, 24, 24, 25, ...
Resampling results across tuning parameters:

          C      RMSE    Rsquared     MAE
 0.25  32.63416  0.7574235  27.93539
 0.50  31.11401  0.7345328  26.30532
 1.00  31.81240  0.7166902  26.97604
 2.00  31.50752  0.7167079  27.08650
 4.00  33.19105  0.7237418  28.28989
 8.00  37.28125  0.7124632  31.47990
16.00  39.46992  0.7136572  33.63473
32.00  39.71876  0.7236006  34.02679
 64.00  38.92351  0.6666740  33.16580
128.00 38.60918  0.6671279  32.86900

Tuning parameter 'sigma' was held constant at a value of 0.07200923
RMSE was used to select the optimal model using the smallest value.
The final values used for the model were sigma = 0.07200923 and C = 0.5.
> summary(svmFitTime)
Length Class Mode
 1   ksvm   S4
```

5<sup>th</sup> Top item

#### Model summary

It shows the model configuration used and some performance measurements.

```
> svmFitTime
Support Vector Machines with Radial Basis Function Kernel

27 samples
 7 predictor

Pre-processing: centered (7), scaled (7)
Resampling: Cross-Validated (10 fold, repeated 5 times)
Summary of sample sizes: 24, 25, 24, 25, 24, 24, ...
Resampling results across tuning parameters:

      C      RMSE    Rsquared     MAE
 0.25  48.82893  0.7129794  38.67265
 0.50  48.75766  0.7402870  39.05445
 1.00  49.78422  0.7329748  40.60900
 2.00  52.34226  0.7338731  44.19451
 4.00  53.97358  0.7348476  46.50983
 8.00  53.97358  0.7348476  46.50983
16.00  53.97358  0.7348476  46.50983
32.00  53.97358  0.7348476  46.50983
 64.00  53.97358  0.7348476  46.50983
128.00 53.97358  0.7348476  46.50983

Tuning parameter 'sigma' was held constant at a value of 1.025851
RMSE was used to select the optimal model using the smallest value.
The final values used for the model were sigma = 1.025851 and C = 0.5.
> summary(svmFitTime)
Length Class Mode
 1   ksvm   S4  ...
```

Linear kernel

1<sup>st</sup> Top item

Model summary

It shows the model configuration used and some performance measurements.

```
| Support Vector Machines with Linear Kernel
```

```
| 27 samples
|   7 predictor
```

```
| Pre-processing: centered (7), scaled (7)
```

```
| Resampling: Cross-Validated (10 fold, repeated 5 times)
```

```
| Summary of sample sizes: 24, 24, 25, 24, 24, 25, ...
```

```
| Resampling results:
```

RMSE	Rsquared	MAE
82.45617	0.6387506	69.15861

```
| Tuning parameter 'C' was held constant at a value of 1
```

```
> summary(svmFitTime)
```

Length	Class	Mode
1	ksvm	S4

3<sup>rd</sup> Top item

Model summary

It shows the model configuration used and some performance measurements.

```
| > svmFitTime
| Support Vector Machines with Linear Kernel
```

```
| 27 samples
|   7 predictor
```

```
| Support Vector Machines with Radial Basis Function Kernel
```

```
| Pre-processing: centered (7), scaled (7)
```

```
| Resampling: Cross-Validated (10 fold, repeated 5 times)
```

```
| Summary of sample sizes: 25, 24, 24, 24, 24, 24, ...
```

```
| Resampling results:
```

RMSE	Rsquared	MAE
40.66145	0.6177963	34.83248

```
| Resampling results across tuning parameters:
```

```
| Tuning parameter 'C' was held constant at a value of 1
```

```
> summary(svmFitTime)
```

Length	Class	Rsquared	MAE
6	Mode	0.7574235	27.93539
1	ksvm	1140	0.7345328
1	S4	0.7345328	26.30532

5<sup>th</sup> Top item

#### Model summary

It shows the model configuration used and some performance measurements.

```
> svmFitTime
Support Vector Machines with Linear Kernel

27 samples
 7 predictor

Pre-processing: centered (7), scaled (7)
Resampling: Cross-Validated (10 fold, repeated 5 times)
Summary of sample sizes: 24, 24, 25, 24, 24, 25, ...
Resampling results:

      RMSE    Rsquared     MAE
 54.3962  0.7154952  43.76852

Tuning parameter 'C' was held constant at a value of 1
> summary(svmFitTime)
Length Class Mode
 1 ksvm S4
```

## Polynomial kernel

1<sup>st</sup> Top item

### Model summary

It shows the model configuration used and some performance measurements.

```
> svmFitTime
Support Vector Machines with Polynomial Kernel

27 samples
 7 predictor

Pre-processing: centered (7), scaled (7)
Resampling: Cross-Validated (10 fold, repeated 5 times)
Summary of sample sizes: 24, 25, 25, 24, 24, 24, ...
Resampling results across tuning parameters:

C      degree  RMSE    Rsquared   MAE
1e-01   1        73.55824  0.6147493  63.76561
1e-01   2        73.46478  0.6151560  63.64232
1e-01   3        73.37252  0.6140579  63.53394
1e-01   4        73.27987  0.6127780  63.42548
1e-01   5        73.18732  0.6107181  63.31717
1e-01   6        73.09354  0.6120832  63.20859
1e+00   1        72.76204  0.6201522  62.89869
1e+00   2        72.10869  0.6332848  62.44691
1e+00   3        71.71738  0.6441611  62.18239
1e+00   4        71.69037  0.6254789  62.27158
1e+00   5        71.62639  0.6219910  62.33192
1e+00   6        71.40476  0.6234699  62.23827
1e+01   1        70.96112  0.6215037  61.90183
1e+01   2        71.99497  0.6187099  62.71736
1e+01   3        73.43456  0.6061303  63.89202
1e+01   4        73.97552  0.6194540  64.33366
1e+01   5        74.38728  0.6210145  64.64864
1e+01   6        74.82973  0.6211369  64.88486
1e+02   1        77.55029  0.5959197  66.56543
1e+02   2        80.61461  0.6062223  67.83912
1e+02   3        82.36920  0.5928625  69.45635
1e+02   4        83.16996  0.5970504  69.97326
1e+02   5        82.74938  0.5727291  69.31515
1e+02   6        81.47318  0.5659279  67.93159
1e+03   1        88.06281  0.5616421  74.49083
1e+03   2        85.39427  0.5585207  71.58462
1e+03   3        75.79175  0.5980777  63.40067
1e+03   4        69.63759  0.6264520  59.42871
1e+03   5        70.17160  0.6510630  60.92572
1e+03   6        73.53432  0.6728755  64.71354
1e+04   1        89.28412  0.5642811  75.48799
1e+04   2        70.33507  0.6487463  61.02457
1e+04   3        85.68248  0.6594849  77.84804
1e+04   4        104.13878 0.6251877  95.07928
1e+04   5        126.59597 0.5931050  113.57099
1e+04   6        151.53478 0.5943769  133.57521
```

Tuning parameter 'scale' was held constant at a value of 0.001  
RMSE was used to select the optimal model using the smallest value.  
The final values used for the model were degree = 4, scale = 0.001 and C = 1000.  
> summary(svmFitTime)  
Length Class Mode  
1 ksvm S4

3<sup>rd</sup> Top item

### Model summary

It shows the model configuration used and some performance measurements.

```
> svmFitTime
Support Vector Machines with Polynomial Kernel

27 samples
 7 predictor

Pre-processing: centered (7), scaled (7)
Resampling: Cross-Validated (10 fold, repeated 5 times)
Summary of sample sizes: 24, 24, 25, 24, 24, 24, ...
Resampling results across tuning parameters:

          C      degree   RMSE    Rsquared   MAE
  1e-01     1       35.19608  0.7208977 30.23723
  1e-01     2       35.15046  0.7209920 30.18742
  1e-01     3       35.10822  0.7210888 30.14117
  1e-01     4       35.07399  0.7188189 30.10235
  1e-01     5       35.04627  0.7183507 30.06829
  1e-01     6       35.01966  0.7222145 30.03816
  1e+00     1       34.94991  0.7126496 29.96176
  1e+00     2       34.65707  0.7050013 29.73213
  1e+00     3       34.24467  0.7152903 29.46053
  1e+00     4       33.86468  0.7202941 29.21025
  1e+00     5       33.69649  0.7178549 29.13168
  1e+00     6       33.66267  0.7139571 29.14174
  1e+01     1       33.45751  0.7229232 29.04526
  1e+01     2       31.12152  0.7282326 26.71747
  1e+01     3       29.81660  0.7387301 25.12709
  1e+01     4       29.25251  0.7536692 24.47846
  1e+01     5       29.05237  0.7603414 24.18877
  1e+01     6       29.21704  0.7561434 24.27790
  1e+02     1       30.41701  0.7306747 25.35022
  1e+02     2       33.41891  0.7006489 28.00263
  1e+02     3       35.00559  0.6856153 29.43440
  1e+02     4       35.93058  0.6812804 30.30483
  1e+02     5       36.40875  0.6800142 30.81734
  1e+02     6       36.61398  0.6815856 31.12266
  1e+03     1       38.37440  0.6649887 32.66915
  1e+03     2       39.10605  0.6646026 33.48628
  1e+03     3       37.35864  0.6854311 31.91964
  1e+03     4       35.09394  0.6951529 29.45974
  1e+03     5       34.54099  0.6844530 28.87964
  1e+03     6       36.71960  0.6968385 30.39085
  1e+04     1       41.01972  0.6568932 34.92762
  1e+04     2       34.70197  0.6851100 29.01608
  1e+04     3       44.66433  0.7105190 35.50840
  1e+04     4       55.19882  0.7152357 42.51253
  1e+04     5       62.22965  0.6932210 48.18695
  1e+04     6       64.53151  0.6920106 50.47946

Tuning parameter 'scale' was held constant at a value of 0.001
RMSE was used to select the optimal model using the smallest value.
The final values used for the model were degree = 5, scale = 0.001 and C = 10.
> summary(svmFitTime)
Length Class Mode
  , , ksvm, S4
```

## 5<sup>th</sup> Top item

### Model summary

It shows the model configuration used and some performance measurements.

```
> svmFitTime
Support Vector Machines with Polynomial Kernel

27 samples
  7 predictor

Pre-processing: centered (7), scaled (7)
Resampling: Cross-Validated (10 fold, repeated 5 times)
Summary of sample sizes: 25, 24, 25, 24, 25, 24, ...
Resampling results across tuning parameters:

          C      degree   RMSE    Rsquared   MAE
  1e-01     1       53.51092  0.6711496  40.43352
  1e-01     2       53.30223  0.6710852  40.24886
  1e-01     3       53.10234  0.6701733  40.06634
  1e-01     4       52.88030  0.6702767  39.86685
  1e-01     5       52.64639  0.6702039  39.67303
  1e-01     6       52.42760  0.6696172  39.48958
  1e+00     1       51.88695  0.6712816  38.88837
  1e+00     2       51.04189  0.6765926  38.05875
  1e+00     3       50.50220  0.6812487  37.55065
  1e+00     4       50.30587  0.6805641  37.47978
  1e+00     5       50.09255  0.6800060  37.42031
  1e+00     6       49.83100  0.6812503  37.18393
  1e+01     1       48.54773  0.6865618  35.89431
  1e+01     2       47.73407  0.6873564  35.43876
  1e+01     3       46.71474  0.6811025  34.75657
  1e+01     4       46.64791  0.6766222  35.00931
  1e+01     5       46.75303  0.6968584  35.27321
  1e+01     6       47.02294  0.7044783  35.78149
  1e+02     1       47.40532  0.7019156  36.59658
  1e+02     2       47.90805  0.6795584  37.98176
  1e+02     3       48.66916  0.6748717  39.02141
  1e+02     4       49.23215  0.6606780  39.55231
  1e+02     5       50.12300  0.6523057  40.33334
  1e+02     6       50.92408  0.6497880  41.05963
  1e+03     1       51.76276  0.6587156  41.55344
  1e+03     2       54.82264  0.6531382  44.27845
  1e+03     3       55.84916  0.6523686  44.88988
  1e+03     4       56.09791  0.6626589  44.69503
  1e+03     5       56.99596  0.6450441  45.48808
  1e+03     6       56.38439  0.6377818  45.41998
  1e+04     1       56.07393  0.6564929  45.19943
  1e+04     2       58.79531  0.6417996  46.99251
  1e+04     3       56.25816  0.6462018  45.44232
  1e+04     4       59.89173  0.6276660  48.44658
  1e+04     5       70.90026  0.5944395  55.64933
  1e+04     6       74.57747  0.5924811  57.92365
```

Tuning parameter 'scale' was held constant at a value of 0.001  
RMSE was used to select the optimal model using the smallest value.

The final values used for the model were degree = 4, scale = 0.001 and C = 10.

```
> summary(svmFitTime)
Length Class Mode
  1  ksvm S4
```

Sigmoid kernel

1<sup>st</sup> Top item

Model summary

It shows the model configuration used.

```
> svmFitTime
Call:
svm(formula = V8 ~ ., data = train, kernek = "sigmoid")
```

Parameters:

```
  SVM-Type: eps-regression
  SVM-Kernel: radial
    cost: 1
    gamma: 0.1428571
  epsilon: 0.1
```

Number of Support Vectors: 27

3<sup>rd</sup> Top item

Model summary

It shows the model configuration used.

```
> svmFitTime
Call:
svm(formula = V8 ~ ., data = train, kernek = "sigmoid")
```

Parameters:

```
  SVM-Type: eps-regression
  SVM-Kernel: radial
    cost: 1
    gamma: 0.1428571
  epsilon: 0.1
```

Number of Support Vectors: 23

5<sup>th</sup> Top item

Model summary

It shows the model configuration used.

```

> svmFitTime

Call:
svm(formula = V8 ~ ., data = train, kernek = "sigmoid")

Parameters:
  SVM-Type: eps-regression
  SVM-Kernel: radial
    cost: 1
   gamma: 0.1428571
  epsilon: 0.1

Number of Support Vectors: 21

```

#### MAPE Summary

In general, the average MAPE of training data sets is much lower than the testing data sets where the ratio for Radial and Sigmoid models is roughly 1:4 and those for Linear and Polynomial are roughly 1:2. It implies Linear and Polynomial models have more stable performance in recognizing the features in data (They have better performance even average MAPE from training and testing sets). It is an interesting finding because they are relatively simple models.

#### Radial kernel

Item	Training	Testing
1st Top Item	11.72%	18.27%
3rd Top Item	18.06%	40.82%
5th Top Item	23.20%	145.20%
Average	17.66%	68.10%

#### Linear kernel

Item	Training	Testing
1st Top Item	14.89%	20.88%
3rd Top Item	19.41%	51.38%
5th Top Item	40.52%	66.68%
Average	24.94%	46.31%

#### Polynomial kernel

Configuration	Item	Training	Testing
method = svmPoly , degree = 4 (Autotune)	1st Top Item	10.73%	18.60%
method = svmPoly , degree = 5 (Autotune)	3rd Top Item	19.46%	45.95%
method = svmPoly , degree = 4 (Autotune)	5th Top Item	39.54%	78.42%
	Average	23.24%	47.66%

Sigmoid kernel

Item	Training	Testing
1st Top Item	10.36%	18.26%
3rd Top Item	13.86%	42.10%
5th Top Item	25.25%	164.32%
Average	16.49%	74.90%

## Holt Winters Analysis

This method applies triple exponential smoothing for level, trend, and seasonal components of time series. It is a predictive forecasting model where the future is predicted by historical data and its trend. Its components can be controlled by parameters. Alpha represents level smoothing; beta specifies trend smoothing and gamma represents seasonality. Seasonality can be classified as additive and multiplicative. In additive seasonality where change is by constant number and multiplicative seasonality changes by a constant factor.

We used tuning grid to optimize the model and select the parameters using metric of MAPE. Also, used ets() function to build the holt winters model.

### MAPE Summary

The averaged MAPE is moderate (34.96%) compared to other models. However, the variations between items are quite high (18.1% to 55.09%). It implies the model may not serve as a generalized method to analyze different products.

Item	Training
1st Top Item	18.10
3rd Top Item	31.69
5th Top Item	55.09
Average	34.96

## ARIMA Analysis

ARIMA stands for AutoRegressive Integrated Moving Average. It is a statistical model for analyzing and forecasting time series data. Below are the key components of this model:

AR (Auto Regression) – It describes the relationship between an observation and some number of past observations.

I (Integrated) – It uses difference between current and one previous time step in order to make it stationary.

MA (Moving Average) – It uses the relationship between an observation and its residual error with moving averages of lagged observations.

These components can be controlled in model by parameters p, d and q respectively. We used a tuning grid to tune these parameters in our modeling.

Arima and Holt winters time series models had similar performance with arima being better with MAPE of 33% while holt winters had 34%. We see from the plots that some outliers have impacted the performance of these models which is expected in these.

### Arima & Holt-Winters Charts

Below are plots of actual data vs prediction from Arima and Holt Winters for three products respectively

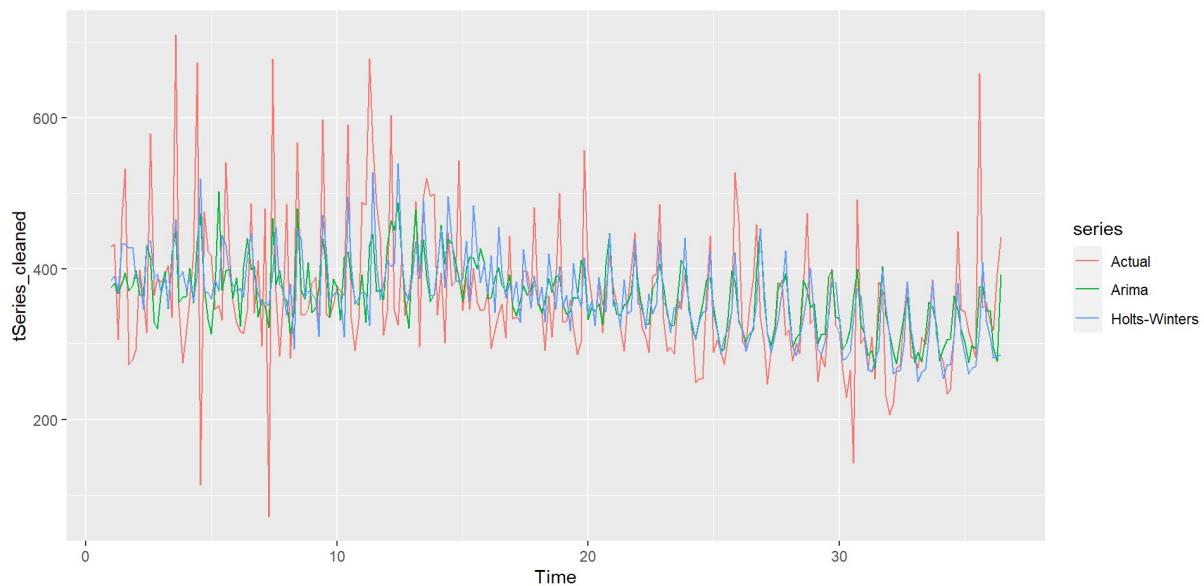


Figure 13

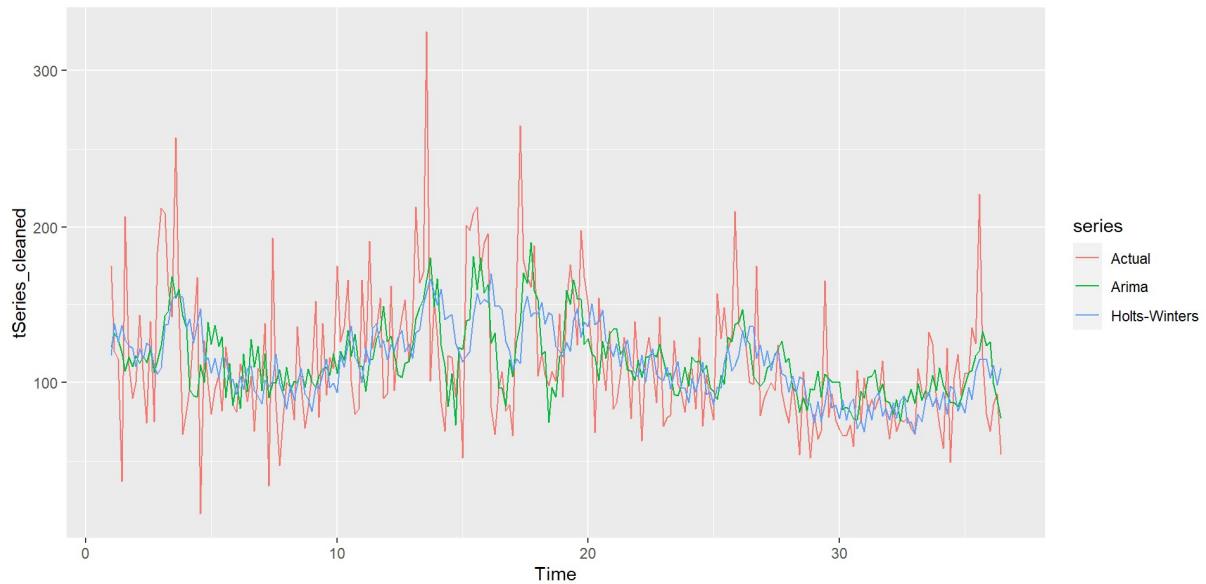


Figure 14

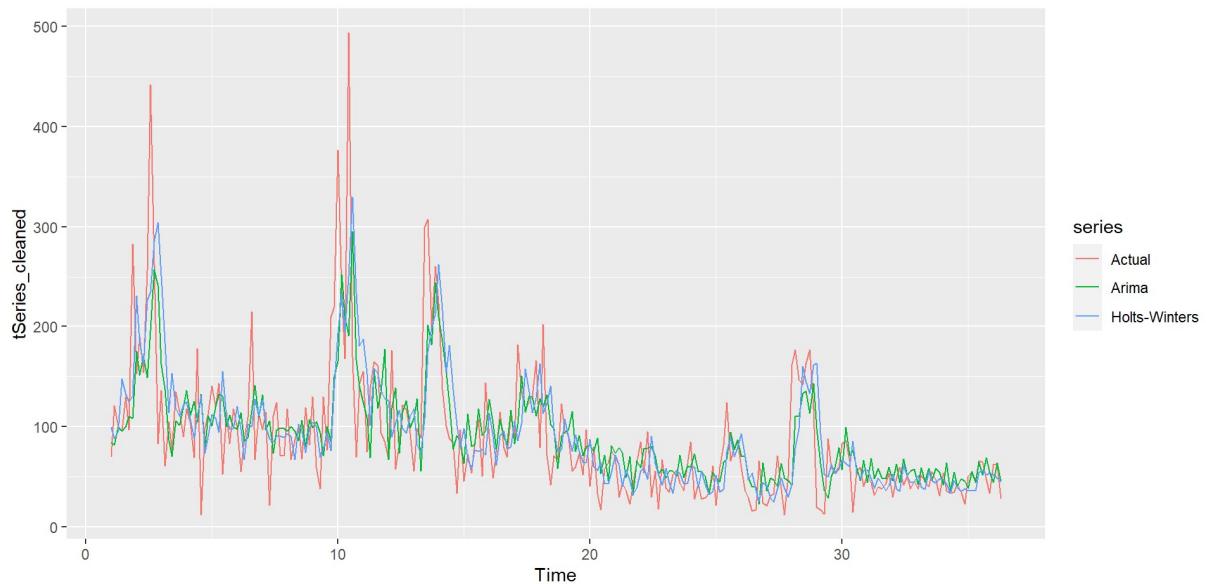


Figure 15

### MAPE Summary

The characteristics is similar to the figures in Holt Winters where it also has moderate level (33.02%) of average MAPE and high variations among products (17.53% to 52.31%). It means it also not suitable to work a generalized solution in prediction.

Item	Training
1st Top Item	17.53%
3rd Top Item	29.21%
5th Top Item	52.31%
Average	33.02%

## Random Forest Analysis

Random forest is a predictive model for classification and regression. It can also be used for time series forecasting after transforming time series into supervised learning problem. For this purpose, we transformed the time series data into 8 columns with 7 independent and one dependent variables. This was just to see how this model works on the time series for assignment purpose only.

Below are plots of actual (red line) vs predicted (blue line) for three products respectively. We see similar kinds of charts in all three products but there are large differences in peaks and valleys which results in large errors.

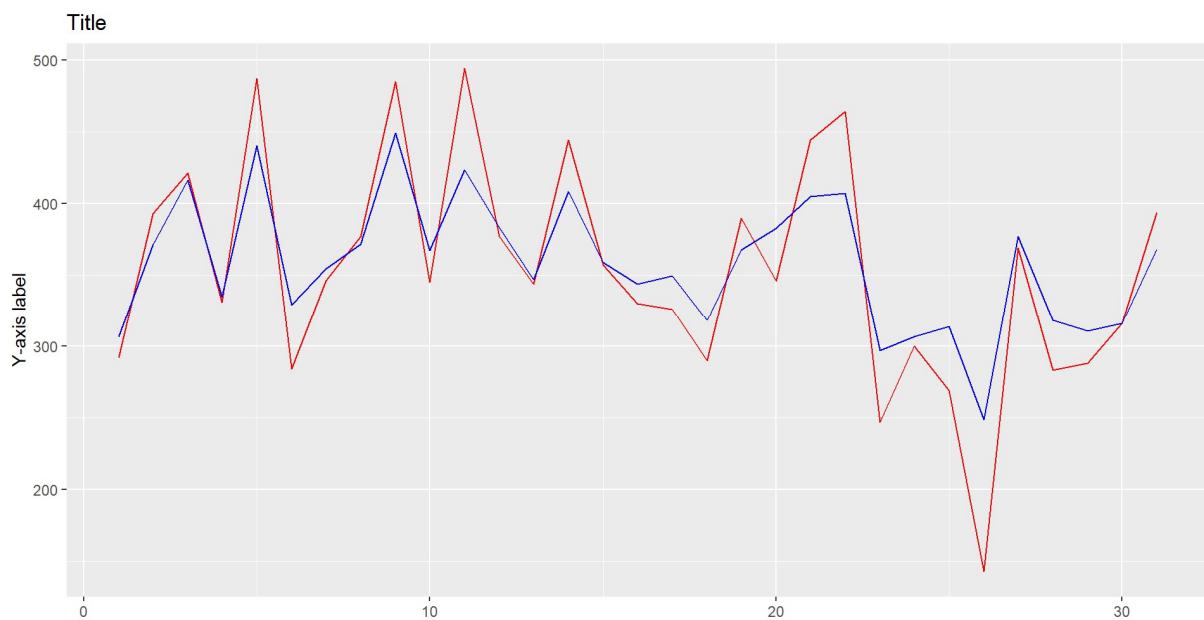


Figure 16

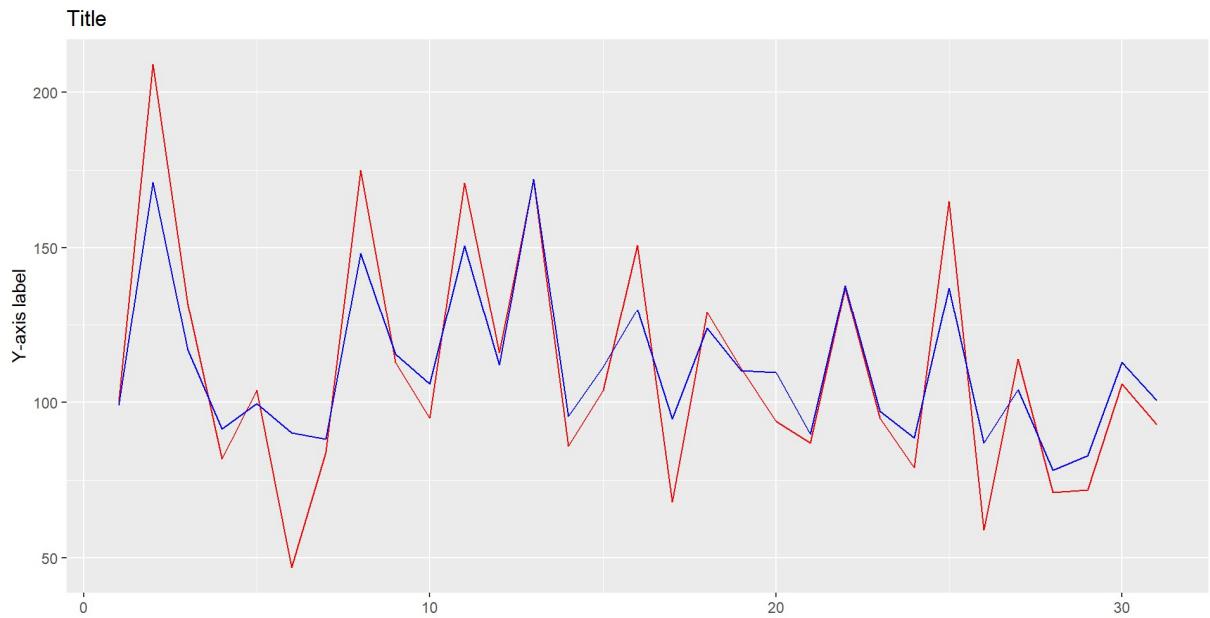


Figure 17

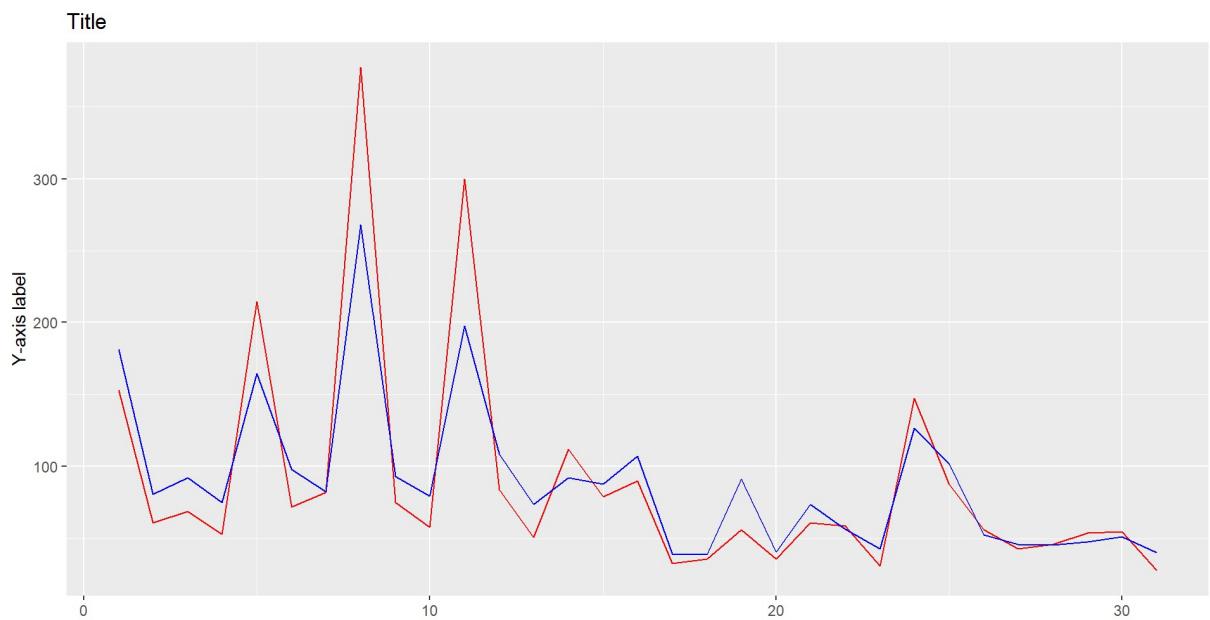


Figure 18

### MAPE Summary

The MAPE are very high (roughly 100 -200 %) it shows Random Forest may not a suitable tools in this prediction.

Item	Training
1st Top Item	93.00 %
3rd Top Item	133.00%
5th Top Item	233.00%
Average	153.00

# **Conclusion**

## **Comparison of all models used**

Overall, all the models did poorly with the best model being averaged neural networks with MAPE of ~30% for training and test set. One hidden layer network produces minimum error ~ 12% in training but in testing it increases to 60% which indicates overfitting.

SVM produced decent models in training, but error increased in testing set. This was true for all the kernel methods. Sigmoid function produced minimum error in training of ~16%.

In time series forecasting models, we found similar performance in both ARIMA and holt winters. Both had an error rate of ~30%.

Linear regression had an error of ~30% in training but in testing set it shot up to 75%. This should be expected since relationships in time series is not linear. Random forest was the worst model with over 150% error.

Analysis Type	Configuration	Item	MAPE (%)	
			Training	Testing
<b>Linear Regression</b>	n/a	1st Top Item	16.91	23.78
		3rd Top Item	23.43	28.53
		5th Top Item	44.80	172.90
		<b>Average</b>	<b>28.38</b>	<b>75.07</b>
<b>Averaged Neural Network</b>	size = 3 (Autotune)	1st Top Item	13.90	24.74
	size = 1 (Autotune)	3rd Top Item	22.40	53.19
	size = 1 (Autotune)	5th Top Item	55.48	25.39
		<b>Average</b>	<b>30.59</b>	<b>34.44</b>
<b>Neural Network</b>	1 hidden layer (15)	1st Top Item	10.35	27.76
		3rd Top Item	9.35	68.94
		5th Top Item	17.18	97.36
		<b>Average</b>	<b>12.29</b>	<b>64.69</b>
	3 hidden layer (5, 5, 5)	1st Top Item	19.91	16.16
		3rd Top Item	21.31	58.21
		5th Top Item	93.81	68.23
		<b>Average</b>	<b>45.01</b>	<b>47.53</b>
<b>SVM</b>	method = svmRadial	1st Top Item	11.72	18.27
		3rd Top Item	18.06	40.82
		5th Top Item	23.20	145.20
		<b>Average</b>	<b>17.66</b>	<b>68.10</b>
	method = svmLinear	1st Top Item	14.89	20.88
		3rd Top Item	19.41	51.38
		5th Top Item	40.52	66.68
		<b>Average</b>	<b>24.94</b>	<b>46.31</b>
	method = svmPoly , degree = 4 (Autotune)	1st Top Item	10.73	18.60
	method = svmPoly , degree = 5 (Autotune)	3rd Top Item	19.46	45.95
	method = svmPoly , degree = 4 (Autotune)	5th Top Item	39.54	78.42
		<b>Average</b>	<b>23.24</b>	<b>47.66</b>
<b>HoltWinters</b>	kernal = "sigmoid"	1st Top Item	10.36	18.26
		3rd Top Item	13.86	42.10
		5th Top Item	25.25	164.32
		<b>Average</b>	<b>16.49</b>	<b>74.90</b>
<b>ARIMA</b>	tuning grid search	1st Top Item	18.10	N/A
		3rd Top Item	31.69	N/A
		5th Top Item	55.09	N/A
		<b>Average</b>	<b>34.96</b>	N/A
<b>Random Forest</b>	tuning grid search	1st Top Item	17.53	N/A
		3rd Top Item	29.21	N/A
		5th Top Item	52.31	N/A
		<b>Average</b>	<b>33.02</b>	N/A
	autotune	1st Top Item	93.00	N/A
		3rd Top Item	133.00	N/A
		5th Top Item	233.00	N/A
		<b>Average</b>	<b>153.00</b>	N/A

## **Recommendation and next step**

We recommend using Averaged neural network and arima models for future prediction of sales quantity. Use this prediction to maintain inventory levels at the store appropriately. Also, keep track of any deviation in the model prediction vs actuals. This divergence signals two things, firstly performance of the product sales is not consistent with past. If it is negative direction then collect more data like feedback from customers and adjust appropriately or offer discounts to boost the sales. If divergence is in a positive direction then continue any sales promotion which are in play. Divergence also means that the model needs to be updated using latest data.

# **Appendix**

## **SQL Script for Data Extraction**

```
select date, item_sk, sum(item_qty) as quantity, sum(item_weight) as weight  
from sales219 where item_sk in ('11740941', '11740923', '11680016', '11610106', '11686823',  
'11741143')  
group by item_sk, date order by date;
```

## **R Script for Data Analysis**

### **Linear Regression Analysis**

```
library(tidyverse)  
library(caret)  
  
setwd("D:/#Spring 2023/5580 - Text Mining/Assignment5")  
  
# setup functions  
mape <- function(actual,pred) {  
  mape <- mean(abs((actual-pred)/actual))*100  
  return (mape)  
}  
  
# Read the data file which has info about top 6 products from sales219  
data <- read.csv('data.csv')  
  
# Get top 6 product ids in list  
top_products <- data %>%  
  group_by(item_sk) %>%  
  summarize(total_quantity = sum(quantity)) %>%  
  select(item_sk, total_quantity) %>%
```

```

arrange(desc(total_quantity))

#-----
# selected product = Top 1st
#11740941
#11741274
#11629829
product <- top_products$item_sk[1]
product

# Filter the data for selected product only
xy <- data %>%
  filter(item_sk == product, quantity != 3) %>%
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well
xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

```

```

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

# Set X and y
X <- train[,1:7]
y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
                             number=10,
                             repeats = 5)

#####
# Linear Regression
# Using pre-sliced data
# Linear regression
glmFitTime <- train(V8 ~ .,
                      data = train,
                      method = "glm",
                      preProc = c("center", "scale"),
                      tuneLength = 10,

```

```

trControl = myCvControl)

glmFitTime
summary(glmFitTime)

# MAPE Training
y_hat = predict(glmFitTime, newdata = X)
mape(y,y_hat)

# MAPE Testing
y_hat2 = predict(glmFitTime, newdata = X_test)
mape(y_test,y_hat2)

#-----
# selected product = Top 3th
#11740941
#11741274
#11629829
product <- top_products$item_sk[3]
product

# Filter the data for selected product only
xy <- data %>%
  filter(item_sk == product, quantity != 1) %>%
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

```

```

# Remove the data as well
xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

# Set X and y
X <- train[,1:7]
y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
  number=10,
  repeats = 5)

```

```

#####
# Linear Regression
# Using pre-sliced data
# Linear regression
glmFitTime <- train(V8 ~ .,
                      data = train,
                      method = "glm",
                      preProc = c("center", "scale"),
                      tuneLength = 10,
                      trControl = myCvControl)

glmFitTime
summary(glmFitTime)

# MAPE Training
y_hat = predict(glmFitTime, newdata = X)
mape(y,y_hat)

# MAPE Testing
y_hat2 = predict(glmFitTime, newdata = X_test)
mape(y_test,y_hat2)

#-----
# selected product = Top 5th
#11740941
#11741274
#11629829
product <- top_products$item_sk[5]
product

```

```

# Filter the data for selected product only

xy <- data %>%
  filter(item_sk == product, quantity != 603) %>%
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well

xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)

rows_to_reshape = total_rows - (total_rows %% 8)

reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)

train<-reshaped_xy[trainIndex,]

test<-reshaped_xy[-trainIndex,]

# Set X and y

X <- train[,1:7]

```

```

y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
                             number=10,
                             repeats = 5)

#####
# Linear Regression
# Using pre-sliced data
# Linear regression
glmFitTime <- train(V8 ~ .,
                      data = train,
                      method = "glm",
                      preProc = c("center", "scale"),
                      tuneLength = 10,
                      trControl = myCvControl)

glmFitTime
summary(glmFitTime)

# MAPE Training
y_hat = predict(glmFitTime, newdata = X)
mape(y,y_hat)

# MAPE Testing
y_hat2 = predict(glmFitTime, newdata = X_test)

```

```
mape(y_test,y_hat2)
```

## Neural Network Analysis

```
library(tidyverse)
```

```
library(caret)
```

```
#setwd("C:/")
```

```
setwd("D:/#Spring 2023/5580 - Text Mining/Assignment5/draft")
```

```
# Read the data file which has info about top 6 products from sales219
```

```
data <- read.csv('data.csv')
```

```
# Get top 6 product ids in list
```

```
top_products <- data %>%
```

```
group_by(item_sk) %>%
```

```
summarize(total_quantity = sum(quantity)) %>%
```

```
select(item_sk, total_quantity) %>%
```

```
arrange(desc(total_quantity))
```

```
# Get the first product and analyze
```

```
#-----
```

```
# selected product = 1st Top
```

```
#11740941
```

```
#11741274
```

```
#11629829
```

```

product <- top_products$item_sk[1]

product

# Filter the data for selected product only
xy <- data %>%
  filter(item_sk == product, quantity != 3) %>% # filter product & outlier
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well
xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

```

```

# Set X and y
X <- train[,1:7]
y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
                             number=10,
                             repeats = 5)

#####
# Neural Network
# Averaged Neural Network

nnFitTime <- train(V8 ~ .,
                     data = train,
                     method = "avNNet",
                     preProc = c("center", "scale"),
                     trControl = myCvControl,
                     tuneLength = 10,
                     linout = T,
                     trace = F,
                     MaxNWts = 10 * (ncol(train) + 1) + 10 + 1,
                     maxit = 500)

nnFitTime

```

```

summary(nnFitTime)

# MAPE Training
y_hat = predict(nnFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)

# Neural Network 1 hidden layer (15)
grid <- expand.grid(layer1 = 15,
                     layer2 = 0,
                     layer3 = 0)

nnFitTime <- train(V8 ~ .,
                     data = train,
                     method = "neuralnet",
                     algorithm="backprop",
                     tuneGrid = grid,
                     learningrate=0.01,
                     preProc = c("center", "scale"),
                     threshold = 0.05,
                     trControl = myCvControl
)

nnFitTime
summary(nnFitTime)

# MAPE Training

```

```

y_hat = predict(nnFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)
plot(nnFitTime$finalModel)

# Neural Network 3 hidden layer (5, 5, 5)
grid <- expand.grid(layer1 = 5,
                     layer2 = 5,
                     layer3 = 5)

nnFitTime <- train(V8 ~ .,
                     data = train,
                     method = "neuralnet",
                     algorithm="backprop",
                     learningrate=0.01,
                     tuneGrid = grid,
                     threshold = 0.05,
                     preProc = c("center", "scale", "nzv"),
                     trControl = myCvControl
)

summary(nnFitTime)

```

```

# MAPE Training
y_hat = predict(nnFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)
plot(nnFitTime$finalModel)

#-----
# selected product = 3rd Top
#11740941
#11741274
#11629829

product <- top_products$item_sk[3]
product

# Filter the data for selected product only
xy <- data %>%
  filter(item_sk == product, quantity != 1) %>% # filter product & outlier
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well
xy$date <- NULL

```

```

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

# Set X and y
X <- train[,1:7]
y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
  number=10,
  repeats = 5)

#####
## Neural Network

```

```

# Averaged Neural Network

nnFitTime <- train(V8 ~ .,
  data = train,
  method = "avNNet",
  preProc = c("center", "scale"),
  trControl = myCvControl,
  tuneLength = 10,
  linout = T,
  trace = F,
  MaxNWts = 10 * (ncol(train) + 1) + 10 + 1,
  maxit = 500)

nnFitTime
summary(nnFitTime)

# MAPE Training
y_hat = predict(nnFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)

# Neural Network 1 hidden layer (15)
grid <- expand.grid(layer1 = 15,
  layer2 = 0,
  layer3 = 0)

nnFitTime <- train(V8 ~ .,

```

```

    data = train,
    method = "neuralnet",
    algorithm="backprop",
    tuneGrid = grid,
    learningrate=0.01,
    preProc = c("center", "scale"),
    threshold = 0.05,
    trControl = myCvControl,
    linear.output = T
)

nnFitTime

summary(nnFitTime)

# MAPE Training
y_hat = predict(nnFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)
plot(nnFitTime$finalModel)

# Neural Network 3 hidden layer (5, 5, 5)
grid <- expand.grid(layer1 = 5,
                     layer2 = 5,
                     layer3 = 5)

```

```

nnFitTime <- train(V8 ~ .,
  data = train,
  method = "neuralnet",
  algorithm="backprop",
  learningrate=0.01,
  tuneGrid = grid,
  threshold = 0.05,
  preProc = c("center", "scale", "nzv"),
  trControl = myCvControl
)

```

```

nnFitTime
summary(nnFitTime)
# MAPE Training
y_hat = predict(nnFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)
# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)
plot(nnFitTime$finalModel)

```

```

#-----
# selected product = 5th Top
#11740941
#11741274
#11629829

```

```

product <- top_products$item_sk[5]
product

# Filter the data for selected product only
xy <- data %>%
  filter(item_sk == product, quantity != 603) %>% # filter product & outlier
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well
xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

```

```

# Set X and y
X <- train[,1:7]
y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
                             number=10,
                             repeats = 5)

#####
# Neural Network
# Averaged Neural Network

nnFitTime <- train(V8 ~.,
                     data = train,
                     method = "avNNet",
                     preProc = c("center", "scale"),
                     trControl = myCvControl,
                     tuneLength = 10,
                     linout = T,
                     trace = F,
                     MaxNWts = 10 * (ncol(train) + 1) + 10 + 1,
                     maxit = 500)

nnFitTime
summary(nnFitTime)

```

```

# MAPE Training
y_hat = predict(nnFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)

# Neural Network 1 hidden layer (15)
grid <- expand.grid(layer1 = 15,
                     layer2 = 0,
                     layer3 = 0)

nnFitTime <- train(V8 ~ .,
                     data = train,
                     method = "neuralnet",
                     algorithm="backprop",
                     tuneGrid = grid,
                     learningrate=0.01,
                     preProc = c("center", "scale"),
                     threshold = 0.1,
                     trControl = myCvControl
)

nnFitTime
summary(nnFitTime)

# MAPE Training
y_hat = predict(nnFitTime, newdata = X)

```

```

mean(100*abs(y_hat-y))/y

# MAPE Testing

y_hat2 = predict(nnFitTime, newdata = X_test)

mean(100*abs(y_hat2-y_test))/y_test)

plot(nnFitTime$finalModel)

# Neural Network 3 hidden layer (5, 5, 5)

grid <- expand.grid(layer1 = 5,
                     layer2 = 5,
                     layer3 = 5)

nnFitTime <- train(V8 ~ .,
                     data = train,
                     method = "neuralnet",
                     algorithm="backprop",
                     learningrate=0.01,
                     tuneGrid = grid,
                     threshold = 0.1,
                     preProc = c("center", "scale", "nzv"),
                     trControl = myCvControl
)

nnFitTime

summary(nnFitTime)

# MAPE Training

y_hat = predict(nnFitTime, newdata = X)

mean(100*abs(y_hat-y))/y

```

```
# MAPE Testing
y_hat2 = predict(nnFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)
plot(nnFitTime$finalModel)
```

## SVM Analysis

```
library(tidyverse)
library(caret)
library(e1071)

setwd("D:/#Spring 2023/5580 - Text Mining/Assignment5")

set.seed(1)
```

```
# Read the data file which has info about top 6 products from sales219
data <- read.csv('data.csv')
```

```
# Get top 6 product ids in list
top_products <- data %>%
  group_by(item_sk) %>%
  summarize(total_quantity = sum(quantity)) %>%
  select(item_sk, total_quantity) %>%
  arrange(desc(total_quantity))
```

```
#-----
# selected product = Top 1st
#11740941
#11741274
```

```

#11629829

product <- top_products$item_sk[1]

product

# Filter the data for selected product only
xy <- data %>%
  filter(item_sk == product, quantity != 3) %>%
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well
xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

```

```

# Set X and y
X <- train[,1:7]
y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
                             number=10,
                             repeats = 5)

# Support Vector Regression
# method = svmRadial
cat(product , "svmRadial")
svmFitTime <- train(V8 ~ .,
                     data = train,
                     method = "svmRadial",
                     preProc = c("center", "scale"),
                     tuneLength = 10,
                     trControl = myCvControl)

svmFitTime
summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y))/y

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test))/y_test

```

```

# method = svmLinear
cat(product , "svmLinear")
svmFitTime <- train(V8 ~ .,
                     data = train,
                     method = "svmLinear",
                     preProc = c("center", "scale"),
                     tuneLength = 10,
                     trControl = myCvControl)

svmFitTime
summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y))/y

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test))/y_test

# method = svmPoly
cat(product , "svmPoly")
svmFitTime <- train(V8 ~ .,
                     data = train,
                     method = "svmPoly",
                     preProc = c("center", "scale"),
                     tuneLength = 10,
                     trControl = myCvControl,
                     tuneGrid = expand.grid(
                     C = c(1e-1, 1e0, 1e1, 1e2, 1e3, 1e4),
                     degree = c(1, 2, 3, 4, 5, 6),

```

```

scale = 0.001)

)

svmFitTime

summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)

mean(100*abs(y_hat-y))/y

y_hat2 = predict(svmFitTime, newdata = X_test)

mean(100*abs(y_hat2-y_test))/y_test

# Kernel = "sigmoid"

cat(product , "sigmoid")

svmFitTime=svm(V8 ~ .,
               data = train,kernek="sigmoid")

svmFitTime

summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)

mean(100*abs(y_hat-y))/y

y_hat2 = predict(svmFitTime, newdata = X_test)

mean(100*abs(y_hat2-y_test))/y_test

#-----
# selected product = 3 rd Top

#11740941

#11741274

#11629829

product <- top_products$item_sk[3]

product

```

```

# Filter the data for selected product only

xy <- data %>%
  filter(item_sk == product, quantity != 1) %>%
  select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well

xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

# Set X and y

X <- train[,1:7]

```

```

y <- train[,8]

X_test <- test[,1:7]
y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
                             number=10,
                             repeats = 5)

# Support Vector Regression
# method = svmRadial
cat(product , "svmRadial")
svmFitTime <- train(V8 ~ .,
                     data = train,
                     method = "svmRadial",
                     preProc = c("center", "scale"),
                     tuneLength = 10,
                     trControl = myCvControl)

svmFitTime
summary(svmFitTime)
y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)
y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)

# method = svmLinear
cat(product , "svmLinear")

```

```

svmFitTime <- train(V8 ~ .,
  data = train,
  method = "svmLinear",
  preProc = c("center", "scale"),
  tuneLength = 10,
  trControl = myCvControl)

svmFitTime
summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)

# method = svmPoly
cat(product , "svmPoly")

svmFitTime <- train(V8 ~ .,
  data = train,
  method = "svmPoly",
  preProc = c("center", "scale"),
  tuneLength = 10,
  trControl = myCvControl,
  tuneGrid = expand.grid(
    C = c(1e-1, 1e0, 1e1, 1e2, 1e3, 1e4),
    degree = c(1, 2, 3, 4, 5, 6),
    scale = 0.001)
)

svmFitTime

```

```

summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y))/y

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test))/y_test

# Kernel = "sigmoid"
cat(product , "sigmoid")

svmFitTime=svm(V8 ~ .,
               data = train,kernek="sigmoid")

svmFitTime

summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y))/y

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test))/y_test

#-----
# selected product = Top 5 th

#11740941
#11741274
#11629829

product <- top_products$item_sk[5]

product

# Filter the data for selected product only
xy <- data %>%

```

```

filter(item_sk == product, quantity != 603) %>%
select(-item_sk)

# Check for any missing dates in the data

is_continuous <- all(diff(xy$Date) == 1)

# Remove the data as well
xy$date <- NULL

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

trainIndex<-createDataPartition(reshaped_xy$V8, p=0.8, list=FALSE)
train<-reshaped_xy[trainIndex,]
test<-reshaped_xy[-trainIndex,]

# Set X and y
X <- train[,1:7]
y <- train[,8]

X_test <- test[,1:7]

```

```

y_test <- test[,8]

myCvControl <- trainControl(method = "repeatedCV",
                            number=10,
                            repeats = 5)

# Support Vector Regression

# method = svmRadial
cat(product , "svmRadial")
svmFitTime <- train(V8 ~ .,
                     data = train,
                     method = "svmRadial",
                     preProc = c("center", "scale"),
                     tuneLength = 10,
                     trControl = myCvControl)

svmFitTime
summary(svmFitTime)
y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y))/y
y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test))/y_test

# method = svmLinear
cat(product , "svmLinear")
svmFitTime <- train(V8 ~ .,
                     data = train,
                     method = "svmLinear",

```

```

preProc = c("center", "scale"),
tuneLength = 10,
trControl = myCvControl)

svmFitTime

summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y))/y

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test))/y_test

# method = svmPoly
cat(product , "svmPoly")

svmFitTime <- train(V8 ~ .,
                     data = train,
                     method = "svmPoly",
                     preProc = c("center", "scale"),
                     tuneLength = 10,
                     trControl = myCvControl,
                     tuneGrid = expand.grid(
                       C = c(1e-1, 1e0, 1e1, 1e2, 1e3, 1e4),
                       degree = c(1, 2, 3, 4, 5, 6),
                       scale = 0.001
                     )
)
svmFitTime

summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y))/y

```

```

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)

# Kernel = "sigmoid"
cat(product , "sigmoid")

svmFitTime=svm(V8 ~ ,
               data = train,kernek="sigmoid")

svmFitTime
summary(svmFitTime)

y_hat = predict(svmFitTime, newdata = X)
mean(100*abs(y_hat-y)/y)

y_hat2 = predict(svmFitTime, newdata = X_test)
mean(100*abs(y_hat2-y_test)/y_test)

```

### **HoltWinters, ARIMA and Random Forest Analysis**

```

library(tidyverse)
library(tseries)
library(forecast)
library(caret)
library(rpart)

# Clear and setup run environment

```

```

setwd('C:/Users/ajay2/OneDrive/OneDrive - Saint Marys University/Desktop/MCDA5580
Mining/Assignmen4')

```

```

# Clear environment
rm(list = ls())

```

```

# Clear plots
while (dev.cur() != 1) dev.off()

# Clear console
cat("\014")

# custom summary function for MAPE
custom_mape <- function(data, lev = NULL, model = NULL)
{
  MAPE <- mean(100 * abs(data$obs - data$pred)/data$obs)
  names(MAPE) <- c('custom_mape')
  return(MAPE)
}

product_fit_function <- function(prod_index)

{

# Get the product data and analyze

product <- top_products$item_sk[prod_index]

# Filter the data for selected product only
xy <- data %>%
  filter(item_sk == product) %>%
  select(-item_sk)

```

```

# Remove the date as well
xy$date <- NULL

# create box plot
ggplot(data = xy, aes(y = quantity)) +
  geom_boxplot(color = 'black', fill = 'white') +
  ggtitle('Box Plot of x')

# Time series

tSeries = ts(xy, start = 1, freq = 7)
plot(stl(tSeries[,1], s.window='periodic'))

# Remove minimum value which is extremely low

if (min(xy) < 10)
{
  print('Removing minimum value from data which is less than 10')
  xy_cleaned <- data.frame(xy$quantity[-which.min(xy$quantity)])
  xy <- xy_cleaned
}

tSeries_cleaned <- ts(xy, start = 1, freq = 7)

# Holt Winters Model with grid search

parameter_grid <- expand.grid(error = c('A', 'M'),

```

```

trend = c('A', 'M'),
seasonal = c('A', 'M'))

# Initialize variables for grid search
best_hw <- NULL
best_mape <- Inf

# Loop through each model combination
for (i in 1:nrow(parameter_grid))
{
  # Extract the model parameters from the grid
  error <- parameter_grid$error[i]
  trend <- parameter_grid$trend[i]
  seasonal <- parameter_grid$seasonal[i]

  model_type = paste0(error, trend, seasonal)

  if (model_type == 'AMA' | model_type == 'MMA' |
    model_type == 'AAM' | model_type == 'AMM')
  {
    next
  }

  # Fit the ETS model
  fit <- ets(tSeries_cleaned, model = model_type)

  # Compute the MAPE
  mape <- mean(100*abs(fitted(fit) - tSeries_cleaned)/tSeries_cleaned)
}

```

```

# Check if this is the best model so far

if (mape < best_mape)
{
  best_hw <- fit
  best_mape <- mape
}

}

# Print the best model and its MAPE value

print(best_hw)

print(paste0('MAPE for holt winters: ', best_mape))

# Use auto tuning arima method and print its coefficients

arima_auto_model <- auto.arima(tSeries_cleaned)

summary(arima_auto_model)

# Plain arima with intuition of weekly repetition of data

arima_plain <- Arima(tSeries_cleaned, order = c(7, 0, 7))

summary(arima_plain)

# Plain arima looks reasonable and simple

print(paste0('MAPE for arima = ',
mean(100*abs(fitted(arima_plain) - tSeries_cleaned)/tSeries_cleaned)))

print(autoplot(tSeries_cleaned, series = 'Actual') +
  autolayer(fitted(arima_plain), series = 'Arima') +
  autolayer(fitted(best_hw), series = 'Holts-Winters'))

```

```

# Random Forest

# Reshape the dataframe to have 7 independent and 1 dependent (8 in total cols)

total_rows = nrow(xy)
rows_to_reshape = total_rows - (total_rows %% 8)
reshaped_xy <-
  as.data.frame(matrix(xy[1:rows_to_reshape, 1], ncol = 8, byrow = TRUE))

# Set X and y

X <- reshaped_xy[,1:7]
y <- reshaped_xy[,8]

# Training Control

myCvControl <- trainControl(method = "repeatedCV",
                             number=10,
                             repeats = 5,
                             summaryFunction = custom_mape)

# Random Forest

random_forest <- train(V8 ~ .,
                        data = reshaped_xy,
                        method = "rf",
                        preProc = c("center", "scale"),
                        tuneLength = 10,
                        metric = 'custom_mape',

```

```

trControl = myCvControl)

y_pred <- predict(random_forest, newdata = reshaped_xy)

print(paste0('MAPE for random forest = ',
round(mean(100 * abs(y_pred - y )/y )), 3))

df <- data.frame(y_actual = y, y_predicted = y_pred)

plot <- ggplot(df) +
  geom_line(aes(x = seq_along(y_actual), y = y_actual), col = "red") +
  geom_line(aes(x = seq_along(y_predicted), y = y_predicted), col = "blue") +
  labs(x = NULL, y = "Y-axis label", title = "Title")

print(plot)
}

```

```
#####

```

```
# Main script starts here
```

```
# Read the data file which has info about top 6 products from sales219
data <- read.csv('data.csv')
```

```
# Get top 6 product ids in list
```

```
top_products <- data %>%
  group_by(item_sk) %>%
  summarize(total_quantity = sum(quantity)) %>%
  select(item_sk, total_quantity) %>%
  arrange(desc(total_quantity))

# Call the function to analyze and fit various models

for (i in seq(1, length(top_products$item_sk), by = 2))
{
  product_fit_function(i)
}
```

## **Reference/ Citation**

Linear Regression

<https://towardsdatascience.com/understanding-linear-regression-output-in-r-7a9cbda948b3>

Neural Network

<https://www.rdocumentation.org/packages/caret/versions/6.0-80/topics/avNNet>

<https://machinelearningmastery.com/how-to-configure-the-number-of-layers-and-nodes-in-a-neural-network/>

<https://machinelearningmastery.com/how-to-control-neural-network-model-capacity-with-nodes-and-layers/>

<https://www.linkedin.com/pulse/choosing-number-hidden-layers-neurons-neural-networks-sachdev>

SVM

<https://www.analyticsvidhya.com/blog/2017/09/understaing-support-vector-machine-example-code/>

<https://bookdown.org/mpfoley1973/supervised-ml/support-vector-machines.html>

<https://dataaspirant.com/svm-kernels/>

<https://koalatea.io/r-svm-regression/>

<https://www.geeksforgeeks.org/introduction-to-support-vector-machines-svm/>

Caret Package

<https://towardsdatascience.com/a-guide-to-using-caret-in-r-71dec0bda208>

<https://cran.r-project.org/web/packages/caret/vignettes/caret.html>

Holt Winters

[Holt-Winters - RapidMiner Documentation](#)

ARIMA

<https://www.investopedia.com/terms/a/autoregressive-integrated-moving-average-arima.asp>

[How to Create an ARIMA Model for Time Series Forecasting in Python - MachineLearningMastery.com](#)

Random Forest for time series forecasting

[Random Forest for Time Series Forecasting - MachineLearningMastery.com](https://machinelearningmastery.com/random-forest-time-series-forecasting/)